

```
In [1]: import torch
import numpy as np
import torchvision
from torchvision import datasets
import torchvision.transforms as transforms
from torchvision.transforms import ToTensor
import torch.nn as nn
import torch.nn.functional as F
import matplotlib.pyplot as plt
```

Loading dataset

```
In [2]: num_workers = 0

batch_size = 64

transform = transforms.ToTensor()

train_data = datasets.MNIST(root='data', train=True,
                             download=True, transform=transform)
test_data = datasets.MNIST(root='data', train=False,
                             download=True, transform=transform)

train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                             num_workers=num_workers)
test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size,
                                             num_workers=num_workers)
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>

0.3%

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to data/MNIST/raw/train-images-idx3-ubyte.gz

100.0%

Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to data/MNIST/raw/train-labels-idx1-ubyte.gz

100.0%

Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz>

2.0%

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to data/MNIST/raw/t10k-images-idx3-ubyte.gz

100.0%

Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz>

100.0%

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to data/MNIST/raw/t10k-labels-idx1-ubyte.gz

Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw

```
In [4]: class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(28*28, 500)

        self.fc2 = nn.Linear(500, 250)

        self.fc3=nn.Linear(250,100)

        self.fc4 = nn.Linear(100, 10)

    def forward(self, x):
        x = x.view(-1, 28*28)
        x = torch.relu(self.fc1(x))

        x = torch.relu(self.fc2(x))

        x= torch.relu(self.fc3(x))
        return F.log_softmax(self.fc4(x), dim=1)

model = Net()
optimizer = torch.optim.Adam(model.parameters(), lr=0.01)#for regularization
                                                    #weight_decay=1e-5

criterion = nn.CrossEntropyLoss()

print(model)
```

```
Net(
  (fc1): Linear(in_features=784, out_features=500, bias=True)
  (fc2): Linear(in_features=500, out_features=250, bias=True)
  (fc3): Linear(in_features=250, out_features=100, bias=True)
  (fc4): Linear(in_features=100, out_features=10, bias=True)
)
```

For training

```
In [5]: n_epochs = 15

model.train() # model for training
trainloss_data=[]

for epoch in range(n_epochs):

    train_loss = 0.0

    for data, target in train_loader:

        optimizer.zero_grad()

        output = model(data)

        #count+=1
        loss = criterion(output, target)

        loss.backward()

        optimizer.step()

        train_loss += loss.item()*data.size(0)

    train_loss = train_loss/len(train_loader.dataset)
```

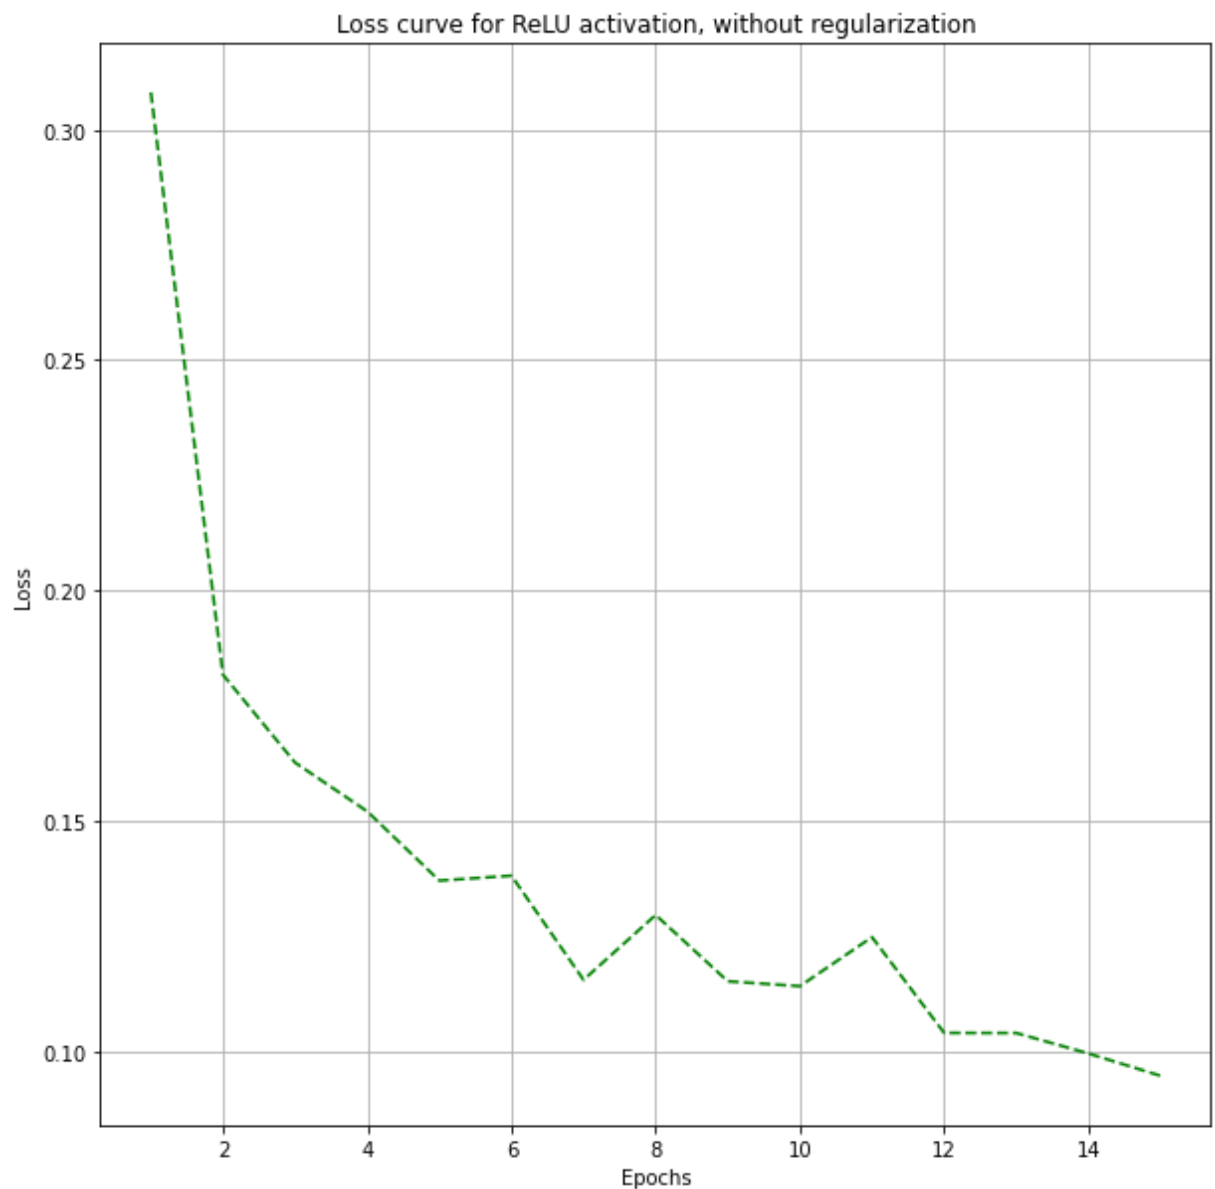
```
trainloss_data.append(train_loss)

print('Epoch: {} \tTraining Loss: {:.6f} '.format(epoch+1,train_loss))
```

```
Epoch: 1      Training Loss: 0.308040
Epoch: 2      Training Loss: 0.181709
Epoch: 3      Training Loss: 0.162660
Epoch: 4      Training Loss: 0.152082
Epoch: 5      Training Loss: 0.137056
Epoch: 6      Training Loss: 0.138126
Epoch: 7      Training Loss: 0.115535
Epoch: 8      Training Loss: 0.129611
Epoch: 9      Training Loss: 0.115225
Epoch: 10     Training Loss: 0.114168
Epoch: 11     Training Loss: 0.124805
Epoch: 12     Training Loss: 0.104027
Epoch: 13     Training Loss: 0.104006
Epoch: 14     Training Loss: 0.099566
Epoch: 15     Training Loss: 0.094724
```

```
In [6]: x=np.arange(1,16,1)
plt.figure(figsize=(10, 10))
plt.plot(x,trainloss_data,'--',color='green')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.grid(True)
plt.title("Loss curve for ReLU activation, without regularization")
```

```
Out[6]: Text(0.5, 1.0, 'Loss curve for ReLU activation, without regularization')
```



For testing

```
In [7]: model.eval()
val_loss, correct = 0, 0
for data, target in test_loader:
    output = model(data)
    val_loss += criterion(output, target).data.item()
    pred = output.data.max(1)[1] # get the index of the max log-probability
    correct += pred.eq(target.data).cpu().sum()

val_loss /= len(test_loader)

accuracy = 100. * correct.to(torch.float32) / len(test_loader.dataset)
print("Accuracy: ", accuracy)
```

Accuracy: tensor(96.5400)

In []: